

preprocess_spy

January 28, 2025

```
[121]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import yfinance as yf
```

```
[122]: df = pd.read_csv('preprocess_spy.csv', header=0, index_col=0)
df.head()
```

```
[122]:
```

	date	dividends	week_day	month	day	quarter	diff_close_close\
0	2020-01-02	0.0	Thursday	1	2	1	2.790588
1	2020-01-03	0.0	Friday	1	3	1	-2.280609
2	2020-01-06	0.0	Monday	1	6	1	1.140289
3	2020-01-07	0.0	Tuesday	1	7	1	-0.843567
4	2020-01-08	0.0	Wednesday	1	8	1	1.594543

	pct_change	pct_open_close	pct_close_open...	ph	spgi	trow	tsla \
0	0.009352	0.005220	0.004111 ...	0.0	0.0	0.0	0.317294
1	-0.007572	-0.011420	0.003892 ...	0.0	0.0	0.0	0.068317
2	0.003815	-0.005955	0.009829 ...	0.0	0.0	0.0	0.056199
3	-0.002811	-0.001915	-0.000898 ...	0.0	0.0	0.0	0.210498
4	0.005329	0.000650	0.004676 ...	0.0	0.0	0.0	0.302827

	tsm	txn	unh	v	wmt	xom
0	0.0	0.0	0.000000	0.0	0.436396	0.000000
1	0.0	0.0	0.000000	0.0	0.241386	0.000000
2	0.0	0.0	0.000000	0.0	-0.410954	0.265997
3	0.0	0.0	0.157431	0.0	0.059462	0.198767
4	0.0	0.0	0.134564	0.0	-0.071081	-0.075503

[5 rows x 53 columns]

```
[123]: df = pd.read_csv('preprocess_spy.csv', header=0, index_col=0)
df['dividends'] = df['dividends'].apply(lambda x: True if x > 0 else False).
    <astype(int)
df.drop(columns=['day', 'quarter', 'seasonal_diff', 'diff_close_close'],
    <inplace=True)
df.drop(columns=['high_open_ratio', 'high_close_ratio', 'low_open_ratio',
```

```
'low_close_ratio'], inplace=True)
```


```
[124]: # from 2020 to 2025
spy = yf.Ticker("SPY")
hist = spy.history(start="2020-01-01", end="2025-01-25")
hist.head()


hist['product_open'] = (hist['High'] - hist['Low']) / (hist['Open'])
hist['product_close'] = (hist['High'] - hist['Low']) / (hist['Close'])

hist['product_diff'] = hist['product_open'] - hist['product_close']
hist['product'] = hist['product_open'] * hist['product_close']

hist['product_diff'].plot()
plt.show();

hist['product'].plot()
plt.show();
```





[]:

[125]: hist.head()

[125]:

	Open	High	Low	Close \
Date				
2020-01-02 00:00:00-05:00	299.961853	301.213476	299.025448	301.194916
2020-01-03 00:00:00-05:00	297.755279	300.054558	297.699654	298.914185
2020-01-06 00:00:00-05:00	297.134101	300.138003	297.013570	300.054565
2020-01-07 00:00:00-05:00	299.479722	299.961845	298.756566	299.210876
2020-01-08 00:00:00-05:00	299.405545	302.038574	299.155231	300.805511

	Volume	Dividends	Stock Splits	Capital Gains\
Date				
2020-01-02 00:00:00-05:00	669151200	0.0	0.0	0.0
2020-01-03 00:00:00-05:00	707709700	0.0	0.0	0.0
2020-01-06 00:00:00-05:00	55653900	0.0	0.0	0.0
2020-01-07 00:00:00-05:00	40496400	0.0	0.0	0.0
2020-01-08 00:00:00-05:00	668296000	0.0	0.0	0.0

	product_open	product_close	product_diff	product
Date				

```

2020-01-02 00:00:00-05:00    0.007294    0.007264    0.000030 0.000053
2020-01-03 00:00:00-05:00    0.007909    0.007878    0.000031 0.000062
2020-01-06 00:00:00-05:00    0.010515    0.010413    0.000102 0.000109
2020-01-07 00:00:00-05:00    0.004025    0.004028   -0.000004 0.000016
2020-01-08 00:00:00-05:00    0.009630    0.009585    0.000045 0.000092

```

```
[126]: hist['trend'] = hist.apply(lambda row: 1 if row['Close'] >= row['Open'] else 0,
    ↪ axis=1)
```

```
[127]: # to datetime
hist['date'] = pd.to_datetime(hist.index)
hist.reset_index(inplace=True, drop=True)

hist['date'] = hist['date'].dt.date
df['date'] = pd.to_datetime(df['date'])
df['date'] = df['date'].dt.date

hist = hist[['date', 'product', 'product_diffrend']]
df.drop(columns=['trend'], inplace=True)
# merge the two dataframes on the date
df = pd.merge(df, hist, on='date', how='left')

df.head(1)
```

```
[127]:      date dividends week_day month pct_change pct_open_close\
0  2020-01-02          0  Thursday    1    0.009352    0.00522

      pct_close_opensecond_diff_seasonalcontinuous_increased\
0          0.004111          3.078827          2

      continuous_decreased.      tsla tsm txn unh  v      wmt xom \
0          0          0 ...  0.317294 0.0 0.0 0.0 0.0 0.436396 0.0

      product product_diffrend
0  0.000053    0.00003    1

[1 rows x 47 columns]
```

```
[128]: # one hot encoding for weekdays

df = pd.get_dummies(df, columns=['week_day'], dtype=int)
```

```
[129]: df.columns
```

```
[129]: Index(['date', 'dividends', 'month', 'pct_change', 'pct_open_close',
      'pct_close_open', 'second_diff_seasonal', 'continuous_increased',
```

```

'continuous_decreased', 'aapl', 'acn', 'adi', 'ame', 'amp', 'amzn',
'aph', 'avgo', 'brk-b', 'cost', 'dov', 'googl', 'hd', 'intu', 'itw',
'jpm', 'lly', 'ma', 'mco', 'meta', 'msft', 'nvda', 'orcl', 'payx', 'ph',
'spgi', 'trow', 'tsla', 'tsm', 'txn', 'unh', 'v', 'wmt', 'xom',
'product', 'product_diff', 'trend', 'week_day_Friday',
'week_day_Monday', 'week_day_Thursday', 'week_day_Tuesday',
'week_day_Wednesday'],
dtype='object')

```

```

[130]: # convert the month integer to datetime and then to string (ex. Jan) and one_
↳hot encode it
df['month'] = df['month'].apply(lambda x: pd.to_datetime(f'2020-{x}-01').
↳strftime('%b'))

df = pd.get_dummies(df, columns=['month'], dtype=int)

df.columns

```

```

[130]: Index(['date', 'dividends', 'pct_change', 'pct_open_close', 'pct_close_open',
'second_diff_seasonal', 'continuous_increased', 'continuous_decreased',
'aapl', 'acn', 'adi', 'ame', 'amp', 'amzn', 'aph', 'avgo', 'brk-b',
'cost', 'dov', 'googl', 'hd', 'intu', 'itw', 'jpm', 'lly', 'ma', 'mco',
'meta', 'msft', 'nvda', 'orcl', 'payx', 'ph', 'spgi', 'trow', 'tsla',
'tsm', 'txn', 'unh', 'v', 'wmt', 'xom', 'product', 'product_diff',
'trend', 'week_day_Friday', 'week_day_Monday', 'week_day_Thursday',
'week_day_Tuesday', 'week_day_Wednesday', 'month_Apr', 'month_Aug',
'month_Dec', 'month_Feb', 'month_Jan', 'month_Jul', 'month_Jun',
'month_Mar', 'month_May', 'month_Nov', 'month_Oct', 'month_Sep'],
dtype='object')

```

```

[131]: df.columns

```

```

[131]: Index(['date', 'dividends', 'pct_change', 'pct_open_close', 'pct_close_open',
'second_diff_seasonal', 'continuous_increased', 'continuous_decreased',
'aapl', 'acn', 'adi', 'ame', 'amp', 'amzn', 'aph', 'avgo', 'brk-b',
'cost', 'dov', 'googl', 'hd', 'intu', 'itw', 'jpm', 'lly', 'ma', 'mco',
'meta', 'msft', 'nvda', 'orcl', 'payx', 'ph', 'spgi', 'trow', 'tsla',
'tsm', 'txn', 'unh', 'v', 'wmt', 'xom', 'product', 'product_diff',
'trend', 'week_day_Friday', 'week_day_Monday', 'week_day_Thursday',
'week_day_Tuesday', 'week_day_Wednesday', 'month_Apr', 'month_Aug',
'month_Dec', 'month_Feb', 'month_Jan', 'month_Jul', 'month_Jun',
'month_Mar', 'month_May', 'month_Nov', 'month_Oct', 'month_Sep'],
dtype='object')

```

```

[ ]:

```

```
[132]: # move trend to the end
```

```
df[['pct_change',  
    'pct_close_open', 'second_diff_seasonal', 'continuous_increased', 'continuous_decreased',  
    'product', 'product_diff', 'aapl', 'acn', 'adi', 'ame', 'amp', 'amzn',  
    'aph', 'avgo', 'brk-b',  
    'cost', 'dov', 'googl', 'hd', 'intu', 'itw', 'jpm', 'lly', 'ma', 'mco',  
    'meta', 'msft', 'nvda', 'orcl', 'payx', 'ph', 'spgi', 'trow', 'tsla',  
    'tsm', 'txn', 'unh', 'v', 'wmt', 'xom']] = df[['pct_change',  
    'pct_close_open', 'second_diff_seasonal', 'continuous_increased', 'continuous_decreased',  
    'product', 'product_diff', 'aapl', 'acn', 'adi', 'ame', 'amp', 'amzn',  
    'aph', 'avgo', 'brk-b',  
    'cost', 'dov', 'googl', 'hd', 'intu', 'itw', 'jpm', 'lly', 'ma', 'mco',  
    'meta', 'msft', 'nvda', 'orcl', 'payx', 'ph', 'spgi', 'trow', 'tsla',  
    'tsm', 'txn', 'unh', 'v', 'wmt', 'xom']].shift(1)
```

```
[133]: df.isna().sum()
```

```
[133]: date                0  
dividends              0  
pct_change            1  
pct_open_close       0  
pct_close_open       1  
..  
month_Mar            0  
month_May            0  
month_Nov            0  
month_Oct            0  
month_Sep            0  
Length: 62, dtype: int64
```

```
[134]: comp_dict = {}  
for i, comp in enumerate(df.columns.tolist()):  
    comp_dict[comp] = i
```

```
[508]:
```

```
[135]: compay_columns = df.columns[comp_dict['aapl']:comp_dict['xom']].tolist()  
  
for comp in compay_columns:  
    df[comp] = (df[comp].shift(-1, fill_value=0) + df[comp])/2
```

```
[136]: # move trend to last column
```

```
trend = df.pop('trend')  
df['trend'] = trend
```


7

```
[138]: trend          1.000000
      pct_open_close  0.059147
      week_day_Monday 0.053844
      brk-b          0.046010
      month_Aug       0.037525
      ...
      orcl           -0.045927
      googl          -0.047263
      week_day_Tuesday -0.055558
      month_Sep       -0.063031
      dividends       -0.064086
      Name: trend, Length: 61, dtype: float64
```

```
[139]: df.describe()
```

```
[139]:
```

	dividends	pct_change	pct_open_close	pct_close_open\
count	1273.000000	1272.000000	1273.000000	1272.000000
mean	0.015711	0.000649	0.000409	0.000235
std	0.124404	0.013184	0.008799	0.009328
min	0.000000	-0.109424	-0.104485	-0.056612
25%	0.000000	-0.005094	-0.002664	-0.004316
50%	0.000000	0.000963	0.000769	0.000747
75%	0.000000	0.007429	0.003962	0.005265
max	1.000000	0.090603	0.060376	0.047994

	second_diff_seasonal	continuous_increased	continuous_decreased
count	1272.000000	1272.000000	1272.000000
mean	0.099614	1.188679	0.790881
std	6.288492	1.566947	1.133974
min	-28.235977	0.000000	0.000000
25%	-3.506115	0.000000	0.000000
50%	0.184875	1.000000	0.000000
75%	3.712875	2.000000	1.000000
max	28.215988	10.000000	7.000000

	aapl	acn	adi ...	month_Feb	month_Jan \
count	1267.000000	1267.000000	1267.000000 ...	1273.000000	1273.000000
mean	0.124624	0.055725	0.022470 ...	0.075412	0.091123
std	0.122300	0.097780	0.072933 ...	0.264159	0.287898
min	-0.285980	-0.416715	-0.237138 ...	0.000000	0.000000
25%	0.052644	0.000000	0.000000 ...	0.000000	0.000000
50%	0.137695	0.000000	0.000000 ...	0.000000	0.000000
75%	0.210602	0.107110	0.000000 ...	0.000000	0.000000
max	0.444473	0.458561	0.446510 ...	1.000000	1.000000

	month_Jul	month_Jun	month_Mar	month_May	month_Nov \
count	1273.000000	1273.000000	1273.000000	1273.000000	1273.000000
mean	0.082482	0.082482	0.087196	0.082482	0.080911
std	0.275206	0.275206	0.282232	0.275206	0.272806
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

	month_Oct	month_Sep	trend
count	1273.000000	1273.000000	1273.000000
mean	0.085625	0.080911	0.552239
std	0.279919	0.272806	0.497459
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	1.000000
75%	0.000000	0.000000	1.000000
max	1.000000	1.000000	1.000000

[8 rows x 61 columns]

```
[140]: # standardize the data
from sklearn.preprocessing import StandardScaler

# columns to standardize

stnd_list = ['pct_change', 'pct_close_open', 'second_diff_seasonal',
             'continuous_increased', 'continuous_decreased', 'product',
             'product_diff', 'aapl', 'acn', 'adi', 'ame', 'amp', 'amzn', 'aph',
             'avgo', 'brk-b', 'cost', 'dov', 'googl', 'hd', 'intu', 'itw', 'jpm',
             'lly', 'ma', 'mco', 'meta', 'msft', 'nvda', 'orcl', 'payx', 'ph',
             'spgi', 'trow', 'tsla', 'tsm', 'txn', 'unh', 'v', 'wmt', 'xom']

scaler = StandardScaler()

df[stnd_list] = scaler.fit_transform(df[stnd_list])

df.dropna(inplace=True)

df.head()
```

```
[140]:      date  dividends  pct_change  pct_open_close  pct_close_open \
1  2020-01-03         0    0.660344    -0.011420      0.415661
2  2020-01-06         0   -0.623821    -0.005955      0.392218
3  2020-01-07         0    0.240200   -0.001915      1.028901
4  2020-01-08         0   -0.262594    0.000650     -0.121477
```

5	2020-01-09	0	0.355112	0.005271	0.476265	
---	------------	---	----------	----------	----------	--

	second_diff_seasonal	continuous_increased	continuous_decreased	aapl \
1	0.473943	0.517975	-0.697716	0.647447
2	-0.453158	-0.758894	0.184486	-0.343627
3	-0.036448	-0.120459	-0.697716	-0.270196
4	-0.493873	-0.758894	0.184486	1.206030
5	0.333805	-0.120459	-0.697716	1.481839

	acn ...	month_Feb	month_Jan	month_Jul	month_Jun	month_Mar \
1	-0.570130 ...	0	1	0	0	0
2	-0.570130 ...	0	1	0	0	0
3	1.207838 ...	0	1	0	0	0
4	1.207838 ...	0	1	0	0	0
5	-0.570130 ...	0	1	0	0	0

	month_May	month_Nov	month_Oct	month_Sep	trend
1	0	0	0	0	1
2	0	0	0	0	1
3	0	0	0	0	0
4	0	0	0	0	1
5	0	0	0	0	1

[5 rows x 62 columns]

```
[176]: #df.drop(columns='date', inplace=True)

from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import accuracy_score

X = df.drop(columns='trend')
y = df['trend']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

clf = DecisionTreeClassifier(max_depth=5)

clf.fit(X_train, y_train)

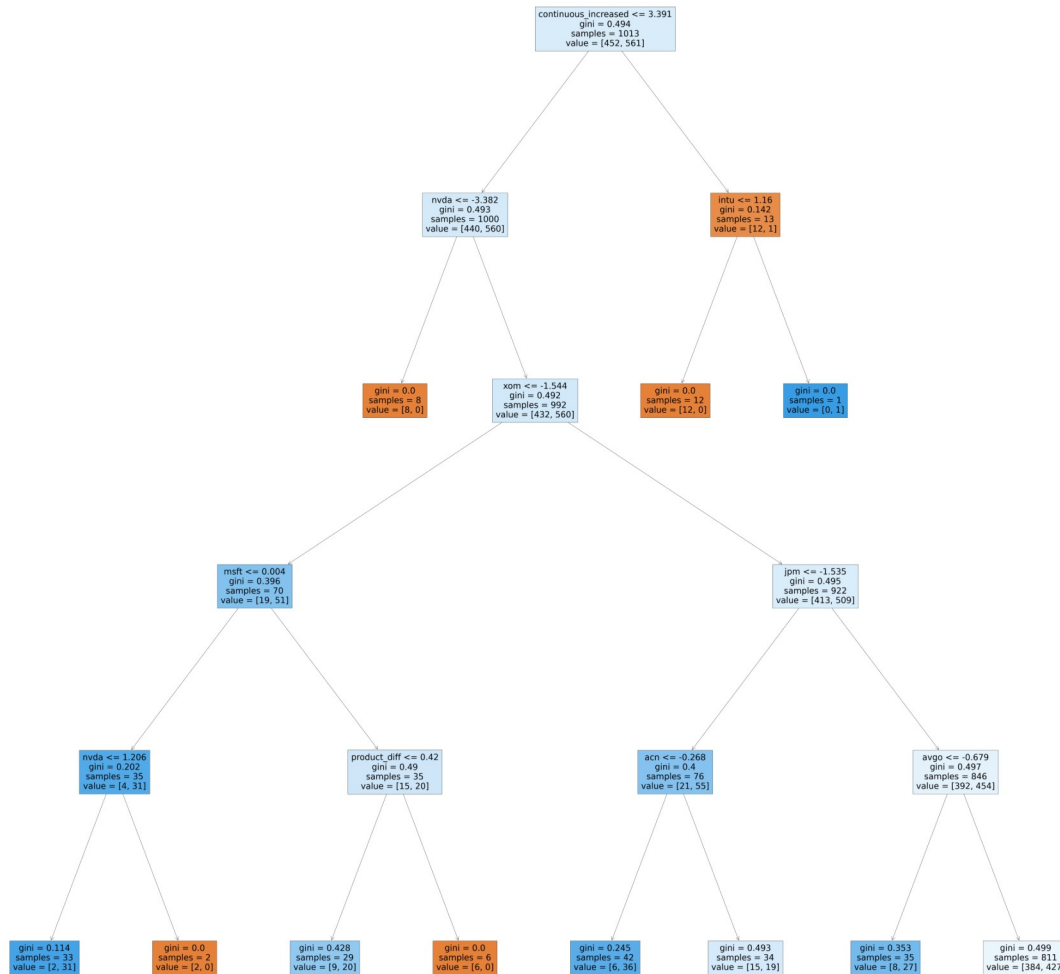
y_pred = clf.predict(X_test)

print(accuracy_score(y_test, y_pred))

plt.figure(figsize=(100, 100))
```

```
plot_tree(clf, filled=True, feature_names=X.columns)
plt.show();
```

0.5354330708661418



```
[173]: feature_importances = clf.feature_importances_
feature_names = X_train.columns
importance_df = pd.DataFrame({'Feature': feature_names, 'Importance':
    feature_importances})
importance_df = importance_df.sort_values(by='Importance', ascending=False)
print(importance_df[:20])
```

Feature Importance

```

4 second_diff_seasonal 0.061119
7 aapl 0.053523
12 amzn 0.044716
2 pct_open_close 0.044683
27 msft 0.044027
39 wmt 0.043474
42 product_diff 0.037222
22 jpm 0.034629
28 nvda 0.032686
3 pct_close_open 0.031751
26 meta 0.028453
33 trow 0.027167
40 xom 0.026803
35 tsm 0.026367
36 txn 0.024482
15 brk-b 0.024224
29 orcl 0.023725
34 tsla 0.023253
23 lly 0.022895
8 acn 0.022194

```

```
[146]: # create a confusion matrix
```

```

from sklearn.metrics import confusion_matrix

confusion_matrix(y_test, y_pred)

```

```
[146]: array([[48, 69],
        [64, 73]])
```

```
[147]: # create a classification report
```

```

from sklearn.metrics import classification_report

print(classification_report(y_test, y_pred))

```

	precision	recall	f1-score	support
0	0.43	0.41	0.42	117
1	0.51	0.53	0.52	137
accuracy			0.48	254
macro avg	0.47	0.47	0.47	254
weighted avg	0.47	0.48	0.48	254

```
[148]: # try xgboost
```

```
from xgboost import XGBClassifier
```

```
xgb = XGBClassifier()
xgb.fit(X_train, y_train)
y_pred = xgb.predict(X_test)
print(accuracy_score(y_test, y_pred))
```

0.4921259842519685

```
[149]: # create a confusion matrix
confusion_matrix(y_test, y_pred)
```

```
[149]: array([[50, 67],
             [62, 75]])
```

```
[170]: # random forest
from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(n_estimators=10)

rf.fit(X_train, y_train)

y_pred = rf.predict(X_test)

print(accuracy_score(y_test, y_pred))

# create a confusion matrix
confusion_matrix(y_test, y_pred)
```

0.5039370078740157

```
[170]: array([[58, 59],
             [67, 70]])
```

```
[182]: # neural network tensorflow

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers

model = keras.Sequential([
    layers.Dense(128, activation='relu', input_shape=[X_train.shape[1]]),
```

```

        layers.Dense(64, activation='relu'),
        layers.Dense(64, activation='relu'),
        layers.Dense(1, activation='sigmoid')
    ])

model.compile(optimizer='adam', loss='binary_crossentropy',
              metrics=['accuracy'])

model.fit(X_train, y_train, epochs=1000)

y_pred = model.predict(X_test)

y_pred = [1 if x > 0.5 else 0 for x in y_pred]

print(accuracy_score(y_test, y_pred))

# create a confusion matrix

confusion_matrix(y_test, y_pred)

```

```

Epoch 1/1000
32/32 [=====] - 1s 980us/step - loss: 0.6968 -
accuracy: 0.5281
Epoch 2/1000
32/32 [=====] - 0s 980us/step - loss: 0.6603 -
accuracy: 0.5953
Epoch 3/1000
32/32 [=====] - 0s 1ms/step - loss: 0.6290 - accuracy:
0.6841
Epoch 4/1000
32/32 [=====] - 0s 850us/step - loss: 0.5884 -
accuracy: 0.7394
Epoch 5/1000
32/32 [=====] - 0s 859us/step - loss: 0.5375 -
accuracy: 0.7572
Epoch 6/1000
32/32 [=====] - 0s 851us/step - loss: 0.4740 -
accuracy: 0.8164
Epoch 7/1000
32/32 [=====] - 0s 790us/step - loss: 0.4041 -
accuracy: 0.8490
Epoch 8/1000
32/32 [=====] - 0s 847us/step - loss: 0.3226 -
accuracy: 0.8894
Epoch 9/1000
32/32 [=====] - 0s 924us/step - loss: 0.2371 -
accuracy: 0.9348
Epoch 10/1000

```